

AURORASTACK LABS

Python AI Engineer (LangChain)

Early-Stage Startup

Full-time

Remote-friendly (US/EU time zones)

AI / Developer Tools

About AuroraStack Labs

AuroraStack Labs is a seed-stage startup building AI-native tools that help teams turn their unstructured knowledge into powerful, conversational applications. We are a small, senior team with backgrounds in applied AI, dev tooling, and product-led growth.

We are looking for a founding Python engineer to lead a new greenfield AI product built on top of LangChain and modern LLM infrastructure.

Role Overview

You will own the design and implementation of our core AI orchestration layer: ingesting data from multiple sources, building retrieval-augmented generation (RAG) pipelines, and exposing these capabilities via robust APIs and internal tools.

This role is ideal for someone who is deeply comfortable with Python, already hands-on with LangChain or similar LLM frameworks, and wants ownership over a 0→1 product in a startup environment.

Python

LangChain

LLMs

RAG

Vector Databases

FastAPI

TypeScript (nice to have)

What You'll Do

- Design, implement, and iterate on AI workflows using Python and LangChain (chains, tools, agents, and custom pipelines).
- Build retrieval-augmented generation (RAG) systems using vector databases (e.g., Pinecone, Weaviate, Qdrant, or similar).

- Integrate with leading LLM providers (OpenAI, Anthropic, etc.) and manage prompt templates, routing, and evaluation.
- Develop production-grade REST APIs and internal services (e.g., with FastAPI) that power our customer-facing applications.
- Implement data ingestion and document processing pipelines (PDFs, HTML, internal tools, knowledge bases, ticketing systems).
- Instrument observability, logging, and evaluation loops to monitor LLM quality, latency, and cost.
- Collaborate closely with design and product to rapidly prototype features and push them to production in short cycles.
- Contribute to technical direction, architecture, and early engineering culture as a founding team member.

Required Experience

- 3+ years of professional software engineering experience, primarily in Python.
- Hands-on experience building applications with LangChain or a similar LLM framework (LlamaIndex, Haystack, Semantic Kernel, etc.).
- Solid grasp of LLM concepts: tokenization, context windows, embeddings, RAG, and prompt design.
- Experience integrating with at least one major cloud provider (AWS, GCP, or Azure) and deploying backend services.
- Strong knowledge of modern Python practices: type hints, virtual environments, packaging, testing, and dependency management.
- Experience building and consuming RESTful APIs; familiarity with frameworks like FastAPI, Flask, or Django.
- Comfortable working in a fast-moving, ambiguous startup environment with high ownership and self-direction.
- Clear written and verbal communication in English; able to explain tradeoffs and decisions to both technical and non-technical stakeholders.

Nice-to-Have Skills

- Experience with one or more vector databases (Pinecone, Qdrant, Weaviate, Chroma).
- Knowledge of basic MLOps practices: experiment tracking, model/version management, monitoring.

- Experience building internal CLIs or developer tools that wrap AI capabilities for other engineers.
- Familiarity with front-end stacks (React, Next.js, or TypeScript) for building basic dashboards or internal UIs.
- Previous experience at an early-stage startup or in a 0→1 product role.
- Exposure to authentication/authorization patterns (OAuth, JWT, SSO).

What We Offer

- Competitive salary plus meaningful early-stage equity.
- Remote-first culture with flexible working hours in US/EU-adjacent time zones.
- Hardware and tooling budget to set up your ideal environment.
- Opportunity to shape a core AI product from the ground up.
- Direct collaboration with the founders and influence over product, culture, and technical roadmap.
- Generous vacation policy and local holidays.

How to Apply

Email jobs@aurorastacklabs.com with the subject "**Python AI Engineer – [Your Name]**" and include:

- Your resume or LinkedIn profile.
- A short note (3–5 paragraphs) on why you are interested in building with LLMs and LangChain at an early-stage startup.
- Links to anything relevant: GitHub, open-source contributions, blog posts, or projects showing AI/LLM work.

If you are unsure whether you meet every requirement but are excited about the role, we encourage you to apply.